

NEW PRODUCT AREA

Education & Re-Skilling

The gap loop tells a person what they are missing. This area tells them how to get it - honestly, demand-first, with the fast route always on the table.

- DESIGN DRAFT
- GROUNDED IN EVIDENCE
- FOURTH MOVE IN THE JOURNEY
- REUSES THE A1 / A2 ARCHITECTURE

The move the master vision misses

The Jobseeker Product Vision is built around presenting *who a person already is* - CV, gaps, interviews. It barely addresses the case where the real blocker is a skills gap that has to be closed. This area fills that hole.

00

THE ONE-LINE SUMMARY

The gap loop tells a person *what they are missing*; this area tells them *how to get it* - honestly, demand-first, with the fast "apply anyway" route always on the table.

Same decode → datafact → submodule → screen architecture as A1 and A2, just pointed at courses instead of CV bullets.

Contents

01 Where it fits

The fourth move, triggered by the honest gap

02 The evidence, and its cautions

What justifies the area, and what constrains it

03 The toolset

Six tools derived from the evidence

04 The screen

How it looks, on the existing architecture

05 Honest-design guardrails

Six rules baked in, from the evidence

06 MVP slice & the one open question


What ships first, and the single build blocker

Where it fits

This is the fourth move in the candidate journey, triggered by the third. The A2 / A3 loop already produces the perfect trigger: the honest outcome where a requirement no reframe can truthfully cover.

01

FIG THE CANDIDATE JOURNEY, AND THE BRANCH THAT WAS A DEAD END



Today the honest "stays a gap" outcome is a dead end. This area turns it into an action path: evaluate whether a course closes it, which course, and whether it is worth it.

THE TRIGGER ALREADY EXISTS

A2 / A3 already produce the honest verdict that a requirement cannot be covered by reframing or adjacent proof. That verdict is the exact signal this area consumes. Nothing new has to be invented to start it - only somewhere for the signal to go.

The evidence, and its cautions

The evidence brief makes re-skilling a first-class need - and forces a set of hard constraints. The numbers are real, but they arrive slowly and only when training is targeted to genuine demand.

02

81-90%

Employed one year after a YH course, with good course-to-job match

SWEDISH YH DATA · §4.4

+7%

Arbetsformedlingen labour-market training vs Komvux - shorter, better aligned

§4.5

2+ yrs

Before training returns appear; near-zero in the short run

HONEST TIMELINE · §4.1

+3%

Komvux vocational income effect, measured at five years

§4.4

What the evidence establishes - and the caution each one forces

Skills gaps are a real, distinct blocker

Separate from CV quality or confidence. For many harder-to-place people, retraining is the actual route back. §4

Returns are real but slow, with lock-in

Training cuts job-search during the course and pays off years later. **The tool must set honest timelines and never over-promise.** §4.1, §8

Targeting to real demand is decisive

Training clears cost-benefit only with strong demand and careful targeting to bottleneck occupations. **Recommend by demand signal, not catalogue browsing.** §4.3, §7.2

Validating can shorten the path

Prior-learning recognition avoids redundant training. **Check "can this be validated?" before recommending a full course.** §4.7

Taxonomies enable the match

ESCO / O*NET give a structured way to map person to course to role. Skill classifiers extract demanded skills from ads, but need expert validation. §4.6

Do not make it digital-first for everyone

Low digital literacy widens inequality. A blended human-plus-guided-digital route is required for some users. §3.4, §5.5

The toolset

Six tools, each one a user need plus the evidence behind it. Tools 1, 3 and 5 are the core MVP; 2 underpins them; 4 and 6 are fast-follows.

03

#	TOOL	USER NEED	WHAT IT DOES	PRIORITY
1	Course-Fit Evaluator	"I have an unbridgeable gap - do I need a course, and which one?"	Takes A2's missing / partial requirements + the person's datafacts + a catalogue + demand signal, and ranks the few courses that actually close the decisive gaps.	CORE MVP
2	Skills-Gap mapper	"What exactly am I missing, in real terms?"	Maps decoded requirements and candidate evidence onto an ESCO / O*NET frame, so gaps are named as skills, not vague asks.	UNDERPINS
3	Demand Signal	"Is this worth retraining for?"	Pulls local and role demand - vacancy data, bottleneck-occupation lists - so recommendations target real demand, not just available courses.	CORE MVP
4	Validering check	"Do I even need the course, or can my experience be recognised?"	Flags where prior learning could be formally recognised - to shorten or skip training - before recommending a full programme.	FAST-FOLLOW
5	Pathway outcomes + honest timeline	"Will this actually pay off, and when?"	Shows the route's evidence (YH employment rate, expected lag) and an honest ramp time - contrasting "fast: apply now" vs "slow: retrain".	CORE MVP
6	Blended / low-friction mode	"I'm not comfortable with digital systems or forms."	Spoken / guided mode plus a coach hand-off for the same flow, so the education path is not gated by digital literacy.	FAST-FOLLOW

THE MVP SLICE


One **course-fit submodule (tools 1 + 2 + 5) reading A2's gaps, rendered on a new screen with the apply-vs-retrain fork and honest timelines. Demand and validering start lightweight; blended mode follows.**

The screen

Reachable from the analysis screen (Ansokningskoll) and from A3 when a gap is marked "stays a gap". The whole point is the fork - people must see the trade-off, not be pushed to train.

04

FIG "CLOSE THIS GAP WITH LEARNING" - THE NEW PANEL, ON THE ANALYSIS SCREEN



The fork is the feature. The evidence says people must see "apply now, lower fit" set beside "retrain, slower, higher fit" and choose - never be nudged onto a course they may not need. Every recommendation cites why it closes the gap, what the outcome and honest timeline are, and offers validering before a full programme.

On the existing architecture

Reuse the HelloLilly framework exactly as A1 / A2 do. No new infrastructure - a cloned submodule, a data lookup, and a new screen.

04

BACKEND SUBMODULE

course-fit

reads: ['meta', 'decodedRole', 'gaps']
writes: ['learningPlan']

Cloned from the decoder scaffold. For each unbridgeable gap, calls the LLM with the gap, the skills frame, the datafacts, the matched courses, and the demand signal, and emits a ranked learningPlan.

THE HONEST GUARDRAIL IS IN THE CONTRACT

Every recommendation must cite *why* - the gap it closes, the demand, the evidence - and must offer the "validate instead of retrain" and "apply now anyway" alternatives. Never a bare "take this course."

The learningPlan shape

```
learningPlan: [{ gapRef, recommendation: 'course' | 'validate' | 'reframe-instead',
  course?, provider, expectedOutcome, rampTime, demandLevel, evidenceNote }]
```

Data sources - the part needing a targeted follow-up

- **Skills frame:** ESCO (EU, Swedish-localised) and / or O*NET.
- **Course catalogue:** YH (MYH), Komvux, SFI, Arbetsformedlingen labour-market training - ideally via their open data / APIs.
- **Demand signal:** vacancy data (the existing LinkedIn job-fetch infra) plus Arbetsformedlingen occupation-outlook (bristyrken).
- **Validering:** MYH / Swedish Council for Higher Education recognition routes - informational first, integrated later.

CROSSLINKING - PER THE VISION'S CENTRAL IDEA

The learning plan crosslinks to coaches with relevant field knowledge, peers who took the same path, course reviews, and the person's Company / Job list. Help appears next to the decision.

Honest-design guardrails

Six rules, baked in from the evidence. They are what keep the area from becoming a course-marketing funnel - which the evidence says would fail the people it is meant to help.

05

THE LOAD-BEARING RULE

Re-skilling is one option, never the default. The fast "apply now anyway" route sits beside "retrain" - and the person decides.

The evidence is blunt: targeting beats volume. A tool that pushes everyone to train fails the people who did not need to.

1 • Never over-promise

Training is slow and period-dependent. Show the lag and the demand caveat. Do not justify it on a 6-12 month horizon. §7.2, §8

2 • Target demand, not catalogue

No recommendation without a demand signal behind it. §4.3

3 • Validate before retrain

Always check prior-learning recognition first. §4.7

4 • Offer the non-training route

Surface "apply now" alongside "retrain". Re-skilling is one option, not the default. §2.5

5 • Do not assume digital comfort

Blended / spoken mode plus coach hand-off for low-digital-literacy users. §5.5

6 • It augments coaches

The robustly effective interventions are human-centred. This tool feeds the coach conversation, it does not replace it. §8

MVP slice & the one open question

The area consumes A2 / A3's gaps, so it ships after the A2 MVP. Everything reuses the framework except one build-blocking unknown.

06

SHIPS FIRST

The MVP

The course-fit submodule (tools 1 + 2 + 5) reading A2's gaps into a learningPlan on the new screen, with the apply-vs-retrain fork and honest timelines. Demand (3) and validating (4) start lightweight / informational; blended mode (6) is a fast-follow.

BUILD ORDER

After the A2 MVP

It depends on A2 / A3 existing, because it consumes gaps. It is the natural next area once the gap loop starts producing unbridgeable gaps - not before.

THE SINGLE OPEN FOLLOW-UP

Do MYH / YH, Komvux, SFI, or Arbetsformedlingen expose an **open course API and an **occupation-demand feed** we can call?**

That is the one build-blocking unknown; everything else reuses the framework. It is worth a short, scoped lookup - not another full deep-research run - before Step 1 of this area.

IN ONE LINE

Same decode → datafact → submodule → screen architecture - just pointed at courses instead of CV bullets, with the fast "apply anyway" route always on the table.